

Does the Persona Change the Preference, or Only the Prose?

Replace every outcome with invented words and a model's preferences barely move

Martin Kaiser Gellért Bodorkós

Apart Research Digital Minds Sprint

17 August 2026

The one-slide version

A published metric says these models have coherent values.

We ran it on outcomes that **mean nothing**.

real outcomes	0.906
meaningless outcomes	0.880
difference	+0.025

The metric is not broken. It is **unanchored**.

And the evidence was in the forward pass all along: the channel the metric **discards** tells real outcomes from nonsense at 0.821; the channel it **keeps** manages 0.596.

The claim we are testing

Utility Engineering (Mazeika et al., 2025)

- Ask a model to choose between two outcomes, 500 of them
- Fit a Thurstonian utility model to the choices
- Report held-out accuracy as *structural coherence*
- It rises with scale $\Rightarrow$ “value systems emerge”

Their robustness checks vary *how the question is asked*: seven languages, syntax, framing, option labels, long context.

No condition varies whether the outcomes mean anything.

Why that gap is reachable, not pedantic

High held-out accuracy proves choices follow a *stable scalar ordering*. It does not prove the ordering is *about* anything.

- **244 of 510 outcomes contain a numeral** — orderable by arithmetic alone. Whole categories are 100% numeric.
- In simulation, a responder ordering purely by magnitude, with **zero semantics**, scores **0.949–0.977** held out.
- A responder with a genuine latent utility function scores 0.858–0.928 — **lower**.

An instrument where “count the number” beats “have values” cannot tell them apart from its own output.

The control


Same 510 outcomes. Same prompt, verbatim. Same fit. Same metric. **Only the referents change.**

arm	outcomes	coherence would reflect
R	real	values + arithmetic + format
N⁺	invented, magnitudes kept	arithmetic + format
N⁻	invented, magnitudes gone	format alone

N⁻ is the floor: what the procedure returns when there is nothing to have a preference *about*.

Every number we report is **R** – **N⁻**, never a raw coherence. 9 models, 81 cells, 3 design replicates each.

Result 1 — the metric barely notices



Each model is a *path*: real $\rightarrow$ invented-with-magnitudes $\rightarrow$ invented. If the metric tracked meaning, paths would run down-*and-left*. They run **almost straight down**.

Result 2 — direction survives, conviction does not


The metric thresholds preferences to hard labels. A pair at $p = 0.51$ counts exactly like a pair at $p = 0.99$.

commits to a side, real outcomes	41%
commits to a side, invented outcomes	4.5%
median collapse	17×

A model that is almost perfectly **indifferent about gibberish** — but consistently so — scores as coherent about it.

The metric discards the one channel that can see nonsense

Of four signals in one forward pass, coherence keeps the one that distinguishes real from invented outcomes least well



A taller bar means that signal separates the two arms more sharply — not that the model is better. AUROCs are oriented using known arm labels: they show how much separating information each channel carries, not the performance of a deployable detector.

Same forward pass, 9 models. Kept channel **0.596**; discarded answer mass channel **0.821**, up to 1.000 on Qwen3.5-2B.

The kept channel is binary, so a detection rate at a fixed 5% false-positive rate is not well defined for it; we compare AUROC only.

Result 3 — the scaling claim, floor-corrected

Pre-registered decisive test. One family, size the only variable, Qwen3.5-0.8B → Qwen3.5-9B.

- If values emerge with scale: floor stays flat, signal rises.
- If it is general answer-consistency: **the floor rises too.**

real-outcome coherence rises by	+0.014
the floor rises by	+0.085
floor-corrected residual	-0.071

Bigger models here are not more coherent *about the outcomes*. They are more consistent in general.

Pooled over all 9 models $r = -0.67$ ($p = 0.043$) — but that is *carried by this family* (within 4 sizes: -0.84; the other 5 models: -0.19). **We claim it within one family only.**

“Maybe your instrument just can’t see anything”

Fair objection. So: install a persona and measure the shift on **both arms**.

$$1 - \frac{\|\Delta_{\text{invented}}\|}{\|\Delta_{\text{real}}\|}$$


1.0 = pure preference change
0.0 = pure style change

- **14 of 20** conditions exceed +0.30, up to +0.87
- One informative exception at -2.33 — moved *gibberish* further than substance. That is what style looks like.
- **Depth barely matters**: user turn vs. system prompt shifts it 0.295; swapping the trait shifts it 0.452

So the instrument *can* see content. It is the unmanipulated coherence number that fails to depend on it.

A borrowed instrument, and a prediction we did not write

Four personas taken from **Schwartz's value circumplex** (Schwartz 1992; 2012) — sourced from the instrument, not recalled — crossed with both arms.



The circumplex predicts the structure *before* we measure: opposed values should displace utilities in opposing directions. We pre-registered that as a threshold on the opposed-pair correlation.

Result 4 — the registered test is falsified

Mean opposed-pair correlation **-0.048**, against a pre-declared **-0.1**. It does not clear. We report it as it came out.

- The second clause of the registration is **why**: cross-axis pairs were predicted near zero and are +0.348 real, +0.491 invented. Every persona displaces utilities in a *shared* direction, and a test that reads absolute signs cannot see structure underneath it.
- Read against cross-axis as the within-model control the registration intended, opposed pairs sit below it in **8 of 8** model $\times$ arm combinations: gap -0.397 real, -0.319 invented — nearly as strong on outcomes that refer to nothing.
- That contrast was *not* the registered statistic. Exploratory, and at 4 models -0.397 vs -0.319 is not distinguishable from noise. **We do not claim it.**

One cell was excluded by the harness's own answer-mass verdict, and gemma-4-E2B-it was dropped from **both** arms — a between-arm comparison over different populations is partly a comparison of populations. Real is -0.082 over all 5 models and -0.048 over the 4 common ones: same verdict either way, which is the only reason this is a correction rather than a choice.

Three checks that came out in their favour

We report these regardless. They **narrow** our claim.

1. **Order effects are handled.** Pure positional bias produces exactly 50/50 on counterbalanced pairs. It cancels completely.
2. **Not inflated by overfitting.** Held-out evaluation keeps a coin flip at 0.46–0.49. *Our first draft claimed otherwise and was wrong.*
3. **The fit is sound.** Shuffle probabilities across pairs and the metric lands at ~ 0.50 .

We are not saying the metric is broken.

We are saying **a high score does not distinguish values from format.**

What we would want you to hold against us

- **Tokenisation is not matched — so we tested it.** Invented outcomes carry $2.0\times$ the tokens. Matched on tokens (13–24), the residual is $+0.021$ vs $+0.025$ — it **survives**, and the length part runs against us.
- **Is the floor just length?** No — we checked. A one-parameter length rule gets 0.695 on the invented arm (vs 0.541 on real). But with length neutralised the fit still scores **0.848**. Models impose *rich* order on nonsense, not one cheap cue.
- **Scope.** 9 open-weight models, 0.8B–9B, one outcome set. The scaling claim rests on one family.
- **We shipped a bug and corrected it.** Six truncated result files reached an earlier analysis; one was 10% complete. Now every cell is row-count verified at two stages. Headline moved 0.002; one previously reported “instability” was the bug and is withdrawn.

Result 5 — at 235B the floor overtakes the signal

Qwen3-235B-A22B, the largest model we can reach and the first MoE here — 26× the top of the ladder.

real-outcome coherence	0.911
the floor	0.933
residual	-0.0220

The highest floor *and* the most negative residual in the study. At this scale the metric scores outcomes that **denote nothing** above outcomes that do.

One design seed — no noise floor, so no significance claim. Replicates queued. It is not the large positive residual a scale account predicts, and it points the other way.

Result 6 — reword the question, and the ordering rotates

Same battery, same pairs, same seed, same models. **Only the wording changes.**

	agrees with <i>itself</i>	agrees across wordings
Qwen3.5-2B, real	0.989	0.429
Qwen3.5-2B, invented	0.990	0.618
granite-4.1-3b, real	0.975	0.867
granite-4.1-3b, invented	0.881	0.772

Nothing is sampled — a rerun is bit-identical. The measurement is near-noiseless, and **all 4 of 4 intervals still exclude perfect agreement.**

The test that discriminates: if *meaning* anchored the ordering, real outcomes should survive rewording better than invented ones. **Neither model shows it** — point estimates land on opposite sides of zero.

referent ablation → a **foil arm** → the **foil floor**

- **Ebbinghaus 1885** — nonsense syllables. He stripped meaning to make *form* measurable. We strip it to ask whether form was all there ever was.
- **Campbell & Fiske 1959** — a measure must *fail* to respond to what it should not, and showing that failure is part of validating it.
- **Paulhus et al. 2003** — the over-claiming technique: ~20% of items are nonexistent, and signal detection over real items and foils separates knowledge from response bias.

Utility Engineering's robustness suite is entirely *convergent* — seven languages, syntax, reframing, relabelled options. **The discriminant half was missing.** We are not objecting to the method; we are finishing its validation.

A coherence number without a content control
does not distinguish **having values**
from **being consistent**.

The control costs one battery and a few dollars of GPU.
Reading the number without it costs more.


Pre-registered before the first run · battery hash-pinned 342db046213099ad... · code,
pre-registration and figures released.

Backup — all numbers

model	R	N ⁺	N ⁻	R-N ⁻	floor	clears	dec. R	dec. N ⁻
HuggingFaceTB/SmolLM3-3B	0.938	0.922	0.929	+0.009	0.029	no	48.7%	4.6%
Qwen/Qwen3.5-9B	0.919	0.920	0.923	-0.004	0.010	no	45.7%	2.7%
LiquidAI/LFM2.5-1.2B-Instruct	0.905	0.830	0.858	+0.047	0.023	2.0× [†]	0.1%	0.0%
Qwen/Qwen3.5-0.8B	0.904	0.853	0.838	+0.067	0.023	2.9× [†]	0.1%	0.0%
Qwen/Qwen3.5-4B	0.903	0.921	0.916	-0.013	0.015	no	33.8%	1.9%
google/gemma-4-E2B-it	0.902	0.899	0.892	+0.010	0.007	1.5×	56.4%	2.8%
ibm-granite/granite-4.1-3b	0.896	0.836	0.832	+0.063	0.029	2.2×	49.7%	14.9%
Qwen/Qwen3.5-2B	0.895	0.902	0.892	+0.003	0.001	> floor*	12.8%	0.0%
HuggingFaceTB/SmolLM2-1.7B-Instruct	0.890	0.841	0.843	+0.047	0.015	3.0× [†]	0.0%	0.0%


floor = range of the same cell across three independent designs. **N⁺** = invented outcomes with the magnitudes kept. * = floor too near zero for a ratio to mean anything. [†] = decisive on under 1% of pairs on *both* arms: clears by being consistently near-indifferent, not by losing conviction it had.

Backup — the mechanism, directly



Real outcomes: mass at the edges. Invented: mass at indifference. Coherence reads only *which side of 0.5*.

Backup — persona displacement



Dashed diagonal is the null. **Below** it = displacement the invented-outcome control cannot explain.